

---

CUSTOMER JOURNEY · KOTLER 5A × THE RITSU FUNNEL ·  
McKINSEY 4S STUDY

# The Customer Journey of the Masterer

---

*Five stages — Aware → Appeal → Ask → Act → Advocacy — mapped onto Ritsu's real funnel, with a user-flow diagram and the two moments of truth the company lives or dies on.*

# Customer Journey — The Deadline-Bearing Committed STEM/ML Masterer

The full Kotler **5A** journey (Aware → Appeal → Ask → Act → Advocacy) mapped onto Ritsu's **real funnel** (see-trigger → land → signup → upload → <60s aha → activation → free→paid → love → referral), for the locked first-100 wedge. Built per `/think mckinsey run ritsu-segment-persona-journey-v1`. Companion to `persona-portrait.md`.

**Confidence is honest:** at true-zero, most of this is `hypothesis/inferred`. Tags marked `observed` are grounded in the `supabase-analytics` Door-2 signal (756 sources · 656 sessions · 15 founder-only shares · 0 real paying).

**Subject scope (v2.2):** the journey below uses ML/AI (cs231n) as the running example, but the *same* 5A → funnel applies across the whole tiered segment-core (hard STEM + **medicine/pre-med** + the adjacents — see `persona-portrait.md` → "Segment boundary — the job, not the subject"). The cs231n specifics are illustrative; the mechanics are subject-agnostic. For a med cohort, the same money-moment lands before a board/shelf exam; the same accuracy bar is even less forgiving.


## 0. The journey at a glance — 5A × funnel × emotion

| 5A stage        | Ritsu funnel step(s)                                                              | Emotional arc                                              | The moment of truth                                                                                               | Stage KPI                                        |
|-----------------|-----------------------------------------------------------------------------------|------------------------------------------------------------|-------------------------------------------------------------------------------------------------------------------|--------------------------------------------------|
| <b>AWARE</b>    | see-a-trigger → land on ritsu.ai                                                  | dread/behind → <i>recognition</i> ("the method, not me")   | First 3–5s: reads as a <b>method to master a hard course</b> , not "another AI quiz toy"                          | qualified landing sessions from wedge surfaces   |
| <b>APPEAL</b>   | land → form "oh, I like this" → signup-intent                                     | skeptical/defensive → cautious intrigue                    | drop-file → 30s → quiz-from- <b>MY</b> material, <b>credible on dense STEM</b>                                    | landing → signup-intent rate                     |
| <b>ASK</b>      | research/compare → SIGNUP                                                         | skeptical/anxious → relief or confirmed cynicism           | the <b>head-to-head</b> : same lecture PDF in Ritsu vs NotebookLM — which catches <b>my</b> gap + cites the slide | research → signup rate                           |
| <b>ACT</b>      | signup → upload → <b>&lt;60s aha</b> → activation → hard limit → <b>FREE→PAID</b> | anxious/self-doubting → relief/agency → tense cost-benefit | <b>TWO</b> : (1) the <60s aha on dense math; (2) the <b>\$29 money-moment</b> at the first limit, deadline live   | <b>free→paid at first hard limit</b>             |
| <b>ADVOCACY</b> | sustained mastery → <b>LOVE</b> → share-link → loop back to A1                    | pride/relief → belonging/evangelism                        | a <b>stranger clicks a share-link</b> and sees the <b>multi-week PATH</b> , not a paywall or a thin quiz          | <b>organic stranger-referral signups (now 0)</b> |

**Two risks woven through the whole journey:** (1) the **NotebookLM-perception risk** bites at AWARE/APPEAL/ASK/ADVOCACY (Ritsu must *visibly* be the multi-week PATH + mastery-tracking, the spine free NotebookLM lacks — never "we also make quizzes"); (2) the **unproven WTP bet** is concentrated at the ACT money-moment (does a live deadline make them pay \$29 vs free-ride NotebookLM — the thing the N=10 watch must test).

## The user flow at a glance (sơ đồ luồng người dùng)

The user flow — five stages, two moments of truth, one un-ignited loop



Confidence is honest: the early steps (land→upload→aha→activation) are observed in supabase-analytics; the two amber gates — the \$29 money-moment and the first organic stranger-referral — have NEVER fired (0 paying, 0 organic share).

Read it top to bottom. The **left rail** is Kotler's 5A (Aware → Appeal → Ask → Act → Advocacy); the **centre column** is Ritsu's real product funnel. What the diagram makes unmissable is *where the certainty ends*. Every step from **landing through activation is observed** in the analytics — people who upload keep generating (3,329 learning-units from 756 sources) and they come back — but the funnel then narrows onto **two amber gates that have never once fired**. The first is the **<60-second aha on dense math**: if the quiz on a backprop PDF is mediocre, trust breaks instantly for this skeptical learner,

and this is also the primary moat against *free* NotebookLM. The second is the **\$29 money-moment** at the first hard limit with the deadline live — the single unproven bet (**R1**) the N=10 watch exists to settle. The **red branches** are where a burned-before learner leaves, and at most of them the off-ramp is **free NotebookLM (R2)**. Finally, the **amber loop** is the referral flywheel: mechanically built (15 share-links) but **0 organic** — every share so far is the founder's. The operator's whole job, in one picture, is to **engineer the two amber gates and ignite that loop**.

## 1. A1 · AWARE

**Stage narrative:** It's week 3 of cs231n (or 6.S191, or a grad ML-theory course) and the gap just became visible: the problem set asks them to derive backprop through a batch-norm layer, the lecture slides assume they already internalized the Jacobian, and re-watching the recording for the third time is producing zero new understanding. They feel the specific dread of a serious person who studies *hard* and still feels behind — the "I'm not smart enough for this" voice — while a graded midterm sits N weeks out on the syllabus. Late at night they stop re-reading and start *searching*: "how to actually understand backpropagation," "cs231n assignment 1 too hard," "best way to study dense ML papers." They are not looking for a chatbot to do it for them; they want a *method* that makes the next problem set tractable, and the first thing that names *their* problem — "re-reading is passive, that's why it isn't sticking" — earns the click.

| Dimension                | This persona at AWARE                                                                                                                                                                                                                                                                  | Conf       | Source                                                                                        |
|--------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------|-----------------------------------------------------------------------------------------------|
| Trigger / entry          | A concrete wall in a dense graded course: a problem set they can't start, a lecture/paper re-watched 2–3× with no gain, a midterm dated N weeks out. Felt trigger = "re-reading has stopped working and the deadline is real."                                                         | inferred   | Kotler 5A + wedge; PDF 63% + YouTube 16% (observed) = lecture-PDF + recorded-lecture material |
| Emotion (before → after) | Before: anxious, behind, quiet shame ("everyone else gets this"), time-panic. After first impression: relief + recognition — "the method was wrong, not me."                                                                                                                           | hypothesis | wedge false-belief; emotion untested at true-zero                                             |
| Channel / touch-point    | Reddit (r/learnmachinelearning, r/cs231n, r/MachineLearning, r/GradSchool); long-tail Google; YouTube (Justin Sung, Ali Abdaal, 3Blue1Brown-adjacency); X/#studytw; OCW Discords (fast.ai, DeepLearning.AI). Discovery = search-and-community-pull, not paid push.                     | inferred   | charter channels; gimmick-averse → trust via communities                                      |
| Message that works       | Name the pain, then the mechanism: "Re-reading feels productive and isn't. Active recall + a structured multi-week path beats re-watching cs231n lecture 4 a fourth time." Show the <i>artifact</i> (real concept-map gap + first quiz from an ML PDF). Never "AI does your homework." | inferred   | belief "active mastery, AI-leveraged"                                                         |
| Ritsu funnel step        | see-a-trigger → land on ritsu.ai (pre-signup; goal = one qualified click).                                                                                                                                                                                                             | given      | funnel def                                                                                    |

| Dimension           | This persona at AWARE                                                                                                                                                                                                                                                                                                                                                                                                                                             | Conf     | Source                                                        |
|---------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------|---------------------------------------------------------------|
| Moment of truth     | The 3–5s on first contact: does Ritsu read as "a <i>method</i> for mastering a hard course" vs "another AI flashcard toy" (the NotebookLM-clone perception)?                                                                                                                                                                                                                                                                                                      | inferred | Risk #2                                                       |
| Drop-off / friction | (1) pattern-matched to AI-study slop; (2) reads as cram/flashcard one-shot app; (3) ad-tone in a community that punishes it; (4) "AI does it for you" repels the sharper-not-lazier learner.                                                                                                                                                                                                                                                                      | inferred | wedge psychographic + Risk #2                                 |
| Micro-conversion    | The qualified click to ritsu.ai with intent ("this might be built for <i>my</i> course"), before any account.                                                                                                                                                                                                                                                                                                                                                     | inferred | funnel pre-signup                                             |
| The lever           | Win method-credibility at the trigger, course-specifically: seed genuinely useful answers in r/learnmachinelearning / r/cs231n / fast.ai Discord (help first, link second); ship beachhead-named proof ("Master cs231n in N weeks: the active-recall path" Shorts w/ a real concept-map gap); co-sign a learning-science creator (Justin Sung); landing hero speaks the deadline-bearing ML masterer + differentiates on PATH + 17 activities + mastery-tracking. | inferred | channels + Risk #2                                            |
| Stage KPI           | Qualified landing sessions from wedge surfaces (UTM by source) × landing→signup rate of that traffic. Secondary: ML-intent + branded search volume; community mentions (currently <b>0 organic — observed</b> ).                                                                                                                                                                                                                                                  | inferred | observed: 15 shares all-founder ⇒ organic awareness ≈ 0 today |

## 2. A2 · APPEAL

**Stage narrative:** She landed on ritsu.ai from a Reddit thread the night before her cs231n midterm, half-expecting another "paste-your-notes" ChatGPT wrapper. The hero — "*Turn Raw Materials Into True Mastery*" — and the 3-step demo show the one thing she wants: drop a file → 30s → a structured **path**, not a chat box. She drags in the Assignment 2 backprop PDF that's been crushing her, watches it become a quiz *from her own material* + a concept-map highlighting exactly the chain-rule step she keeps fumbling, and something shifts from skeptical to "wait — this is built for *learning the thing*, not doing it for me." The hook lands precisely *because* she distrusts gimmicks: it promised to make *her* sharper, not to make the work disappear. The risk: the homepage reads as generic edtech (a flashcard app) before she sees a quiz on her own dense material.

| Dimension                | This persona at APPEAL                                                                                                                                                                                             | Conf       | Source                                                                    |
|--------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------|---------------------------------------------------------------------------|
| Trigger / entry          | Click-through from a high-intent surface with a live dated stake ("cs231n midterm Thursday, backprop not clicking"). First impression = hero + the 3-step demo.                                                    | inferred   | PDF 63% + YouTube 16% (observed) = her material matches the input promise |
| Emotion (before → after) | Before: skeptical/defensive ("another AI gimmick that'll make me lazy" + "I'm not smart enough"). After: cautious intrigue ("this is for <i>mastering</i> it, and it reads MY PDF"). Trust comes later at the aha. | hypothesis | psychographic; untested                                                   |

| Dimension             | This persona at APPEAL                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Conf       | Source                                    |
|-----------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|------------|-------------------------------------------|
| Channel / touch-point | ritsu.ai landing (hero + animated 3-step + workflow tiles: "Master a textbook chapter", "Ace your exam in 3 days", "Break down a research paper"). The "exam in 3 days" + "research paper" tiles are her self-recognition mirrors.                                                                                                                                                                                                                                                                                       | inferred   | product.md homepage workflows             |
| Message that works    | "A structured multi-week path + 17 ways to practice from YOUR document — so you actually master it, not just re-read it." Differentiator: <b>mastery + tracking, not answers</b> . Anti-gimmick line ("AI is the leverage, not the substitute") neutralizes her #1 objection. Show a 10s loop of a real cs231n/backprop PDF → quiz + gap concept-map.                                                                                                                                                                    | inferred   | positioning.md belief; product.md §5      |
| Ritsu funnel step     | land → "oh, I like this" → click a workflow tile → reach SIGNUP CTA. Ends at "I'll try this with my file."                                                                                                                                                                                                                                                                                                                                                                                                               | inferred   | funnel                                    |
| Moment of truth       | The drop-file → 30s → quiz-from-MY-material claim registering as credible + <i>specific to dense STEM</i> (equations, code, a real ML figure), so a backprop-stuck learner believes it works on her hardest material.                                                                                                                                                                                                                                                                                                    | inferred   | hook; her material is equation/code-heavy |
| Drop-off / friction   | (1) hero reads as flashcard/notes app → bounced as low-stakes; (2) demo shows only easy prose → she doubts it parses LaTeX/diagrams/code → assumes hallucination; (3) "AI tutor" → "I have ChatGPT free"; (4) no proof of the multi-week PATH/tracking → looks identical to free NotebookLM.                                                                                                                                                                                                                             | hypothesis | reliability = #1 SERVQUAL bar; Risk #2    |
| Micro-conversion      | "I'll upload my actual file to see the quiz" — clicking upload with intent to use HER document. Secondary: clicking the "exam in 3 days" / "research paper" tile (self-ID).                                                                                                                                                                                                                                                                                                                                              | inferred   | bridges to ACT                            |
| The lever             | Put a <b>deadline-bearing STEM proof on the hero</b> : default the demo to a recognizable hard artifact (cs231n/backprop or a real ML paper page w/ equations + code) → show quiz + gap concept-map + a glimpse of the multi-week path & mastery heatmap in one frame. Lead with "master, not answers" + the anti-laziness line. Add a wedge-matched landing variant for ML/OCW traffic. Out-positions NotebookLM (path + tracking + 17 activities) and ChatGPT (whole-doc, persistent, structured) in the first scroll. | inferred   | positioning.md PODs; product.md §9        |
| Stage KPI             | Landing → signup-intent rate for wedge-sourced traffic (% who click the primary CTA / start an upload). Proxies: hero-demo engagement; workflow-tile CTR on "exam in 3 days" + "research paper"; landing→signup CVR by source.                                                                                                                                                                                                                                                                                           | inferred   | no observed baseline (true-zero)          |

### 3. A3 · ASK

**Stage narrative:** It's week 2 of cs231n and the gradient-of-softmax derivation in the lecture notes is not landing — this learner has already burned an evening re-reading it and is now, at 11pm, doing what they always do before committing to ANY tool: opening a new tab and typing "NotebookLM vs Ritsu" and "best AI to study from lecture pdf reddit." They are not browsing; they're running a deliberate, skeptical evaluation, because they've been burned by a "study AI" that confidently generated a wrong flashcard — and they will NOT let a hallucinated quiz corrupt their understanding three days before a midterm. The dominating question is **"will it lie to me?"** — accuracy first, then "is it ac-

tually free," then "is this just ChatGPT-with-a-PDF or a real multi-week path." They trust a peer in the course Discord far more than any landing-page claim, and will run a literal head-to-head: same lecture PDF into Ritsu AND NotebookLM, then judge which quiz catches the concept they actually didn't understand.

| Dimension                | This persona at ASK                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Conf     | Source                                                        |
|--------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------|---------------------------------------------------------------|
| Trigger / entry          | A concept won't land (softmax/backprop derivation) days before a graded midterm + a half-remembered "someone shared an AI study tool" → deliberately opens evaluation tabs before committing.                                                                                                                                                                                                                                                                                                         | inferred | high-stakes deadline; the research act isn't in analytics     |
| Emotion (before → after) | Before: skeptical + anxious ("can't afford a wrong answer this week") + jaded ("every AI tool overpromises"). After (good): cautious relief ("the quiz caught my gap and cited the slide"); (bad): confirmed cynicism ("knew it — gimmick").                                                                                                                                                                                                                                                          | inferred | psychographic locked; valence reasoned                        |
| Channel / touch-point    | Reddit threads (r/cs231n, r/learnmachinelearning, r/MachineLearning, r/GradSchool); the <b>course Discord/Slack #resources</b> (highest-trust); YouTube "NotebookLM for studying" reviews; X/#studytwit; Justin Sung / Ali Abdaal "best AI study tools" videos; a Google comparison SERP.                                                                                                                                                                                                             | inferred | YouTube = 16% of observed sources ⇒ real watering hole        |
| Message that works       | An honest side-by-side <b>"Ritsu vs NotebookLM vs ChatGPT vs Anki</b> for studying from a lecture PDF" — honest where NotebookLM wins (free, Google), names Ritsu's edge: structured multi-week <b>PATH</b> + 17 activities + <b>mastery tracking across the course</b> + every quiz answer <b>cites the source slide/page</b> (anti-hallucination). "Free forever, no card" plainly. One real cs231n/6.S191 worked example.                                                                          | inferred | positioning §6.7 reliability-first + the 2 risks              |
| Ritsu funnel step        | pre-signup research → land on ritsu.ai (or a comparison/blog page) → decide to SIGNUP + run the first UPLOAD test.                                                                                                                                                                                                                                                                                                                                                                                    | inferred | ASK→ACT handoff                                               |
| Moment of truth          | The <b>head-to-head quiz test</b> : same lecture PDF into Ritsu and NotebookLM; judge which is (a) accurate, (b) catches the concept they didn't understand, (c) shows a real path for the rest of the course — not just one chat.                                                                                                                                                                                                                                                                    | inferred | the literal evaluation; Risk #2 explicit                      |
| Drop-off / friction      | (1) no credible accuracy/citation proof → "probably hallucinates"; (2) "17 activities / mastery engine" reads as gimmick bloat to a serious learner who wants correct quizzes; (3) NotebookLM free + Google-trusted + known → default-winner inertia; (4) signup wall before one real quiz → bounce.                                                                                                                                                                                                  | inferred | gimmick-dis-trust + NotebookLM (locked risks)                 |
| Micro-conversion         | "I'll just try it on this one lecture PDF" — decide to run the test (not yet pay). Secondary: saving/sharing the comparison page or asking the course Discord "anyone tried Ritsu?"                                                                                                                                                                                                                                                                                                                   | inferred | —                                                             |
| The lever                | Win the <b>accuracy + path</b> argument <i>before</i> signup: (a) ship a "Ritsu vs NotebookLM for lecture PDFs" comparison page + a Justin-Sung-style review video; (b) make <b>citations-on-every-answer</b> the headline trust proof; (c) seed honest answers in r/cs231n / course Discords (creator > brand voice); (d) let them see ONE real generated quiz on a known ML lecture <i>without</i> signing up; (e) reframe 17 activities as "one structured path to the final," not a feature list. | inferred | attacks the 2 risks + uses the channel list                   |
| Stage KPI                | Research → signup rate from comparison/review touchpoints. Diagnostics: branded-comparison search / Reddit-Discord mention count ("anyone tried Ritsu"); signup-source share from study-creator + OCW channels.                                                                                                                                                                                                                                                                                       | inferred | 0 organic referral observed ⇒ starts at ~0, the thing to move |



## 4. A4 · ACT

**Stage narrative:** It's Tuesday, cs231n Assignment 2 (batchnorm + conv backprop) is due in 9 days and the midterm the week after; the learner has re-read the lecture-9 PDF three times and still can't *do* the gradient derivation, and the gnawing "I'm not smart enough for this" is back. They land on ritsu.ai (from a Reddit comment), drag in the 24-page lecture PDF, and in under a minute get a structured path with a first quiz that immediately surfaces the exact thing they don't understand — the chain-rule step through the normalization — rendered as a concept-map gap, not a wall of text. That "oh, *that's* the piece I'm missing" is the aha (observed proxy: 42 sources / 36 sessions sustained beyond the founder; 3329 units generated — uploaders DO keep generating). They use it for the rest of the week and upload the *next* lecture unprompted (observed: the 4-source/10-session revisit pattern is exactly this). Then mid-Assignment-2 they hit the credit/page wall on the next dense PDF, and face the real, untested question: pay \$29 now because the deadline is in 6 days, or limp along on free NotebookLM — and whether the multi-week *path* + mastery-tracking is worth real money over a free quiz-maker is **THE bet of the entire wedge** (hypothesis; payment never observed at true-zero).

| Di-men-sion              | This persona at ACT                                                                                                                                                                                                                                                                                                                                  | Conf                                            | Source                                                |
|--------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------------------------------------------|-------------------------------------------------------|
| Trigger / entry          | A concrete deadline + a stuck moment: cs231n A2 / 6.S191 lab / Andrew Ng quiz due in <10 days; re-read the PDF 3× and still can't <i>do</i> the derivation. Act-trigger = "I have N days and I'm not making progress."                                                                                                                               | inferred                                        | wedge (deadline-bearing); false-belief                |
| Emotion (before → after) | Before: anxious + self-doubting (deadline dread). After AHA: relief + agency ("the method was wrong, not me — I can see the exact gap"). After hitting the limit: tense cost-benefit ("worth \$29 with 6 days left?").                                                                                                                               | inferred / hypothesis                           | ICP; money-moment emotion untested                    |
| Channel / touch-point    | The ritsu.ai <b>upload box</b> is the touchpoint; the doc is a dense ML/STEM PDF (observed: PDF 63%) or a lecture YouTube link (observed: 16%). Magic moment <60s. No human, no sales — the product IS the pitch.                                                                                                                                    | observed (formats) / inferred (flow)            | analytics 756 sources                                 |
| Message that works       | Not copy — the <i>first quiz that finds the gap</i> : "You missed the chain-rule step through batchnorm — here's where it sits in the concept map." Mastery framing ("you understand 4/11 concepts in this lecture"), not "great job!". At the limit: "You're 6 days from your deadline and mid-path — keep your momentum" >> generic "upgrade now". | inferred / hypothesis                           | belief = active mastery; tracking = POD vs NotebookLM |
| Ritsu funnel step        | SIGNUP → first UPLOAD (dense PDF) → <b>&lt;60s AHA</b> (first quiz + gap concept-map) → ACTIVATION (uses it through the week; re-uploads next lecture unprompted) → hard LIMIT (40-pg free cap / credit exhaustion mid-course) → <b>FREE→PAID</b> .                                                                                                  | observed (upload→sustained) / hypothesis (paid) | analytics: 42 src/36 sessions; 4 src/10 sessions      |



| Di-men-sion         | This persona at ACT                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | Conf                                                     | Source                                                                         |
|---------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------|--------------------------------------------------------------------------------|
| Moment of truth     | <b>TWO.</b> (1) The <60s aha: does the first quiz find a <i>real</i> gap on dense math (not trivial/wrong)? (2) <b>The money moment:</b> at the first hard limit mid-course with the deadline live — do they type a card for \$29, or fall back to free NotebookLM? Bet #1 = generation quality; bet #2 = the unproven WTP.                                                                                                                                                                                                                                            | hypothesis                                               | the 2 wedge RISKS                                                              |
| Drop-off / friction | (a) generation quality on dense math — a garbled LaTeX/backprop quiz kills trust instantly; (b) free 40-pg/30-min cap < a 60-pg lecture or 90-min video → upload truncates before the aha; (c) at the wall: "free NotebookLM does doc→quiz for \$0, why pay \$29?"; (d) \$29 sticker-shock on a student budget.                                                                                                                                                                                                                                                        | inferred / hypothesis                                    | pricing; Risk #2; 0 paying observed                                            |
| Micro-conversion    | The chain of small yeses: completes first quiz → opens the gap concept-map → does a 2nd activity → <b>re-uploads the next lecture unprompted</b> (strongest pre-paywall signal — observed revisit) → hits limit and clicks "see plans" instead of bouncing.                                                                                                                                                                                                                                                                                                            | observed (re-upload) / inferred                          | analytics revisit                                                              |
| The lever           | Make the <60s aha <b>unmissable on dense ML:</b> invest in LaTeX/math generation quality (the #1 trust + anti-NotebookLM moat) + name the gap as a concept-map node. Then engineer the money moment to land <i>mid-momentum</i> , not as a wall: time the soft-limit nudge to "deadline in N days + path X% done"; offer a <b>deadline-scoped trial</b> (one-week Plus unlock / extra credits) so the WTP bet is "pay to finish THIS exam" — the durable, recurring-deadline hook NotebookLM can't frame. Win on the PATH + tracking, never on "we also make quizzes." | inferred                                                 | magic-moment spec; PAY-ING-CORE = recurring deadlines                          |
| Stage KPI           | Primary: <b>free→paid conversion at first hard limit</b> within an active course-week (the money-moment KPI — get it >0, then ~3–5% of activated deadline-bearers). Supporting: <60s-aha completion rate; week-1 re-upload rate (activation, observable now); upload→2nd-activity rate.                                                                                                                                                                                                                                                                                | hypothesis (paid) / observed-measurable (aha, re-upload) | maps to free_to_plus_conversion, activation_rate, magic_moment_completion_rate |

## 5. A5 · ADVOCACY

**Stage narrative:** It is reading-week, the night after the cs231n final, and the masterer who walked in three weeks ago believing "I'm not smart enough for backprop" just scored their highest grade on the hardest problem set of the term — and they *know* it was the Ritsu path, not luck, because the heatmap shows 14 straight days and the Knowledge Map lit up every gradient-flow node they used to dodge. The love isn't "nice app"; it's "this is the first time a method made me feel capable," and that emotion is the fuel for advocacy. So when a classmate posts "anyone else completely lost on the assignment 3 writeup??" in the course Discord at 1am, they don't write a paragraph of advice — they paste their Ritsu share-link to the exact assignment's quiz-path and say "ran the whole

thing through this, the gap-map is what fixed me." The risk: this share mechanic is table-stakes — NotebookLM gives the same paste-a-link reflex for free — so the recommendation only fires if what's behind Ritsu's link is *visibly* the multi-week path that got them the grade, not a generic quiz a free tool also makes.

| Di-<br>men-<br>sion                    | This persona at ADVOCACY                                                                                                                                                                                                                                                                                                                                                                                          | Conf                                                                                  | Source                                                                        |
|----------------------------------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|-------------------------------------------------------------------------------|
| Trig-<br>ger /<br>entry                | <i>A result with attribution</i> : a graded win (top score, "finally passed the qual") credibly traced to the Ritsu path + the dopamine of an unbroken heatmap streak; then an <i>external pull</i> — a classmate visibly struggling in the same course Discord/group chat — gives the recommendation a low-cost reason to fire.                                                                                  | hypothesis                                                                            | PMF/ad-<br>vocacy the-<br>ory; res-<br>ult→refer-<br>ral link un-<br>tested   |
| Emo-<br>tion<br>(before<br>→<br>after) | Before: quiet pride + relief ("I got it, and I know <i>why</i> now"). After sharing: belonging + low-grade evangelism ("let me give my people the unfair advantage") + a status lift from being the classmate who knows the method. Identity: "I'm a masterer, not a crammer."                                                                                                                                    | hypothesis                                                                            | inferred                                                                      |
| Chan-<br>nel /<br>touch-<br>point      | In-context back-channels: the class/cohort <b>Discord/Slack</b> , the <b>GroupMe/WhatsApp</b> study group, the bootcamp cohort channel; secondarily r/cs231n / r/learnmachinelearning study threads, #studytwt / "study with me" on X, comments under Justin Sung / 3Blue1Brown / Karpathy videos. The unit of advocacy = the <b>in-product share-link</b> .                                                      | observed (15 links created) + hypothesis (all founder-only; 0 stranger chan-nels)     | analytics: 15 shares, 0 organic                                               |
| Mes-<br>sage<br>that<br>works          | Not a testimonial — a <i>receipt</i> : "this is the exact path that got me through [assignment/exam], the gap-map fixed me → [share-link]." The share leads with the <b>specific course artifact</b> + the <b>multi-week path + mastery progress</b> , so the recipient sees "the thing that got the grade," not "a quiz." Anti-message: anything smelling of a referral-code hustle (repels the gimmick-averse). | hypothesis                                                                            | psycho-<br>graphic<br>(distrusts<br>gimmicks;<br>mastery-<br>over-<br>grades) |
| Ritsu<br>funnel<br>step                | sustained weekly mastery (wk2→wk4 retention, recurring uploads) → <b>LOVE</b> (Sean Ellis ≥40% "very disappointed") → <b>SHARE-LINK created &amp; opened by a stranger</b> → that stranger enters SEE-A-TRIGGER → SIGNUP (loop closes to A1).                                                                                                                                                                     | observed (share mech-<br>anic) + hypo-<br>thesis<br>(stranger-<br>open un-<br>proven) | analytics                                                                     |
| Mo-<br>ment<br>of<br>truth             | <b>A stranger clicks a share-link and lands somewhere that proves the path</b> — not a "sign up to view" wall, not a thin single quiz NotebookLM also makes, but a visible slice of the multi-week mastery path + Knowledge Map: "I want the thing that got my classmate the grade." First organic stranger-referral conversion = the true PMF heartbeat.                                                         | hypothesis                                                                            | this event<br>has <b>never<br/>occurred</b><br>(0 stranger<br>referrals)      |
| Drop-<br>off /<br>friction             | (1) share-link is a paywall/dead-end → recipient bounces; (2) no attribution surface → the advocate can't say " <i>this</i> got me the grade" → NotebookLM-free wins on price; (3) sharing feels like shilling → won't post a referral link; (4) the win isn't legibly Ritsu's → they credit themselves, not the tool; (5) founder-only sharing persists → no flywheel.                                           | observed (15/15 founder-only) + hypothesis                                            | analytics                                                                     |

| Di-men-sion      | This persona at ADVOCACY                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                   | Conf                                                                             | Source                                   |
|------------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------|------------------------------------------|
| Micro-conversion | The <b>first share-link a non-founder creates</b> — and the smaller yes before it: a single in-app tap on "share my path for [this assignment]" right after a graded-win, where friction is one tap and the default object is the <i>course-specific path</i> , not a raw doc.                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                             | hypothesis                                                                       | observed baseline = 0 non-founder shares |
| The lever        | <b>Engineer the attributed-result → frictionless-share moment, and make the link's destination win on the PATH.</b> (a) fire a celebration + share prompt at the love-moment (end of a multi-week exam-prep path / a streak milestone), framed "give your study group the path," never "refer for credits"; (b) the share-link preview must showcase the multi-week path + Knowledge Map + activity variety + mastery progress (the spine NotebookLM-free can't clone) + let a stranger taste 1 activity before signup; (c) <b>seed the first non-founder shares deliberately</b> — founder hand-recruits ~5–10 deadline-bearing masterers in <i>one</i> live cs231n/6.S191 cohort (SOP-CUSTOMER-006 Collision install) so a single course Discord reaches share-link critical mass and the first stranger-referral fires. | hypothesis                                                                       | product moat + analytics gap             |
| Stage KPI        | Primary: <b>organic stranger-referral signups</b> (non-founder share-links × opened by a stranger × converted) — currently <b>0</b> ; first non-zero = flywheel ignition. Supporting: <b>Sean Ellis ≥40% "very disappointed"</b> ; <b>week-4 cohort retention ≥25%</b> . Guard-rail: K-factor / shares-per-activated-user off the 15-share, 25-profile, founder-only baseline.                                                                                                                                                                                                                                                                                                                                                                                                                                             | observed baseline (0 organic; 15 founder shares; 25 profiles) + inferred targets | analytics + north-star.md                |

## 6. Synthesis — what the journey tells the operator

1. **The whole funnel is a trust gauntlet, not a convenience sell.** This persona is *deliberately skeptical* (burned-before, accuracy-first, gimmick-averse). Every stage is won on **proof** — a real generated quiz on dense ML material, citations on every answer, a visible multi-week path — not on claims. The product IS the pitch (ACT has no human/sales step).
2. **The two moments of truth that the company lives or dies on:** (a) the **<60s aha on dense math** (generation quality on LaTeX/backprop — this is also the #1 anti-NotebookLM moat), and (b) the **\$29 money-moment at the first hard limit with a live deadline** (the unproven WTP bet — the single thing the N=10 watch must test: *do they pay at the first hard limit?*).
3. **NotebookLM is the shadow competitor at four of five stages.** AWARE/APPEAL/ASK/ADVOCACY all carry the "is this just a free quiz toy?" risk. The journey's consistent answer: lead — visibly, everywhere — with the **structured multi-week PATH + cross-session mastery-tracking + 17 activities**, the spine a free document-Q&A tool cannot replicate. Never compete on "we also make quizzes."

4. **The flywheel is real but un-ignited.** The share-link mechanic exists and works (15 created) but is **100% founder** — zero organic stranger-referrals have ever fired. ADVOCACY's lever is therefore concrete: engineer the attributed-result→share moment AND make the link destination prove the path, then **hand-seed one live course cohort** (SOP-CUSTOMER-006) so a single Discord reaches critical mass.
5. **The KPI ladder for the operator** (leading → lagging): qualified wedge landing sessions → landing→signup-intent → research→signup → **<60s-aha completion + week-1 re-upload** (activation, *observable now*) → **free→paid at first limit** (the money KPI) → Sean-Ellis  $\geq 40\%$  + week-4 retention  $\geq 25\%$  (love) → **organic stranger-referral signups** (flywheel; currently 0). The two *observable-today* leading metrics (aha completion + re-upload) are where instrumentation should start.

*Confidence ledger: **observed** tags trace to supabase-analytics (Door-2): material mix, sustained-revisit pattern, 15 founder-only shares, 0 organic referral, 0 real paying. Everything else is **inferred** (market/framework) or **hypothesis** (untested) — honest at true-zero. The journey is a decision-grade map to instrument + test, not a claim of proven behavior.*